



Homework 4 - CS361 Threaded Merge Sort

Freddie Patterson

Table Data and Plots:

Array Size	Merge	TMerge 2	TMerge 4	TMerge16	TMerge32
1	0	0	0	0	0
2	0	252	0	0	0
4	0	357	322	0	0
8	1	350	264	1	0
16	2	226	389	1202	1
32	4	229	307	625	1469
64	8	153	375	758	1174
128	12	169	233	605	1613
256	26	188	281	639	1039
512	76	262	232	715	1121
1024	108	220	243	664	1248
2048	319	350	311	704	1239
4096	599	634	451	736	1241
8192	1174	1052	750	1034	1803
16384	2386	1609	1358	1865	2153
32768	5098	3229	1979	3251	3194
65536	11292	6087	4181	4870	5243
131072	22497	12562	7800	8422	9441
262144	46267	23917	14812	14157	15650
524288	91362	48723	29517	25143	29092
1048576	180591	111573	62070	42700	52344





Conclusion on data:

The data clearly shows that as thread count increases, the overall time it takes to complete the merge sort decreases. This is not the case necessarily at the low end of the size of the array. I would postulate that this is to do with the fact that the time it takes to even assign the threads, takes longer than it actually takes to sort the array. However, as the value of n items in the array increases, it becomes very clear that the time taken greatly decreases with the number of threads. This evidence would show that the time complexity of the regular merge sort of is $O(n\log(n))$, and although it keeps generally to the improvement that would be expected in $O(n/t\log(n/t)) + O(n)$, the data clearly shows inconsistency between the improvement by threads that does not follow it exactly, even if it follows the same trend generally e.g. The differences between 131072 and 262144 items.

What I have found easiest has to have been getting the overall structure in of the program, not just for the standard single threaded merge sort, but getting the structure for running the threads simultaneously and making sure that things that need to be locked and performed differently from the single threaded version of the merge sort, were surprisingly easy to set up the structure of the program.

The hardest has definitely been getting the rejoining of all the threads into the correct merged list. Although not too difficult to do with the lesser number of threads, when running at 16 or 32 threads, the merging of the whole list became very difficult and something I struggled quite a lot with. Also just merge itself. I made the mistake of trying to just directly implement from the written version of the merge function, of which I made quite a lot of syntax mistakes that took me a decent amount of time to debug especially with so many threads running. E.g., declaring an `int* list` without giving it a size value, and then accidentally making it just an `int` of that value instead of a size.